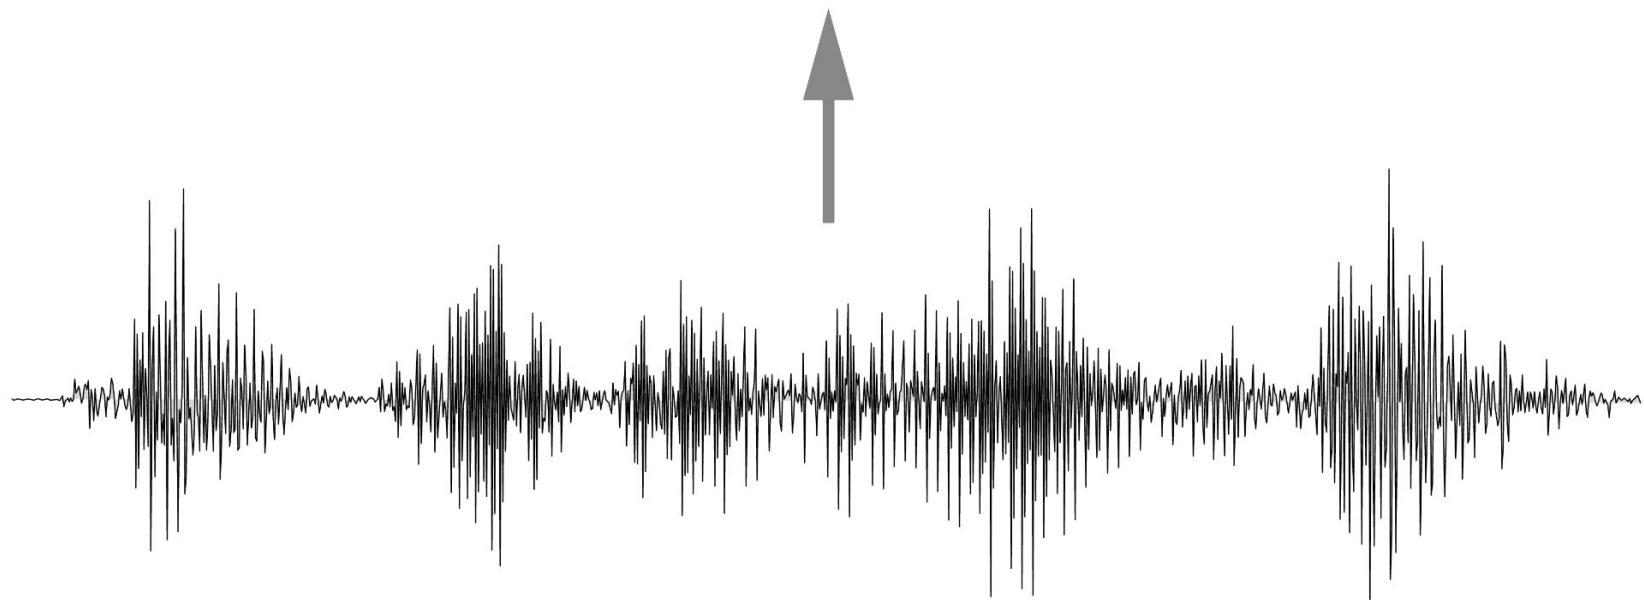


CTC



j u m p s o v e r t h e l a z y d o g



Speech recognition: The input can be a spectrogram or some other frequency based feature extractor.

Connectionist Temporal Classification (CTC): is a way to get **around not knowing the alignment** between the input and the output. As we'll see, it's especially well suited to applications in speech recognition.

Mapping input sequences such as audio, to corresponding output sequences

$$X = [x_1, x_2, \dots, x_T]$$

$$Y = [y_1, y_2, \dots, y_U]$$

There are challenges which get in the way of us using simpler supervised learning:

- Both X and Y can vary in length.
- The ratio of the lengths of X and Y can vary.
- We don't have an accurate alignment (correspondence of the elements) of X and Y .

Overcomes these challenges

The **CTC algorithm overcomes** these challenges.

For a given X it gives us an output distribution over all possible Y 's.

We can use this distribution either to *infer* a likely output or to assess the *probability* of a given output.

Loss Function

We'd like to train our model to **maximize** the probability it assigns to the right answer.

We'll need to efficiently compute the conditional probability $p(Y|X)$.

The function $p(Y|X)$ should also be differentiable, so we can use gradient descent.

Inference

Naturally, after we've trained the model,
we want to use it to infer a likely Y given an X . This means solving

$$Y^* = \operatorname{argmax}_Y p(Y \mid X).$$

Algorithm

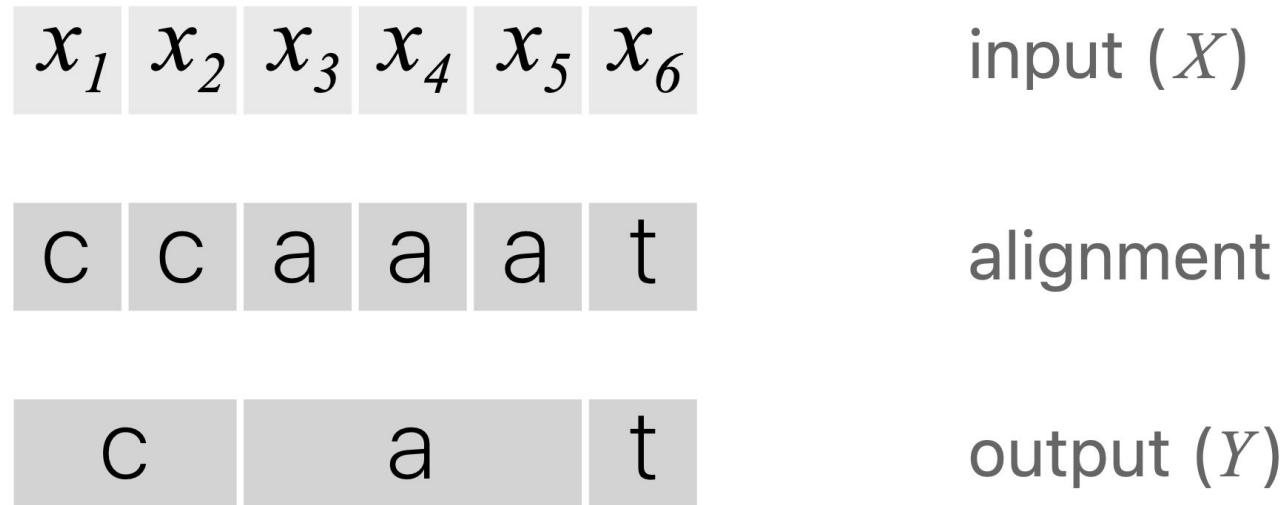
The CTC algorithm is *alignment-free*

it doesn't require an alignment between the input and the output.

However, to get the probability of an output given an input, CTC works by summing over the probability of all possible **alignments** between the two.

Alignment

First consider a naive approach. Let's use an example. Assume the input has length six and $Y = [c, a, t]$. One way to align X and Y is to assign an output character to each input step and collapse repeats.



This approach has two **problems**.

- It doesn't make sense to force every input step to align to some output.
- We have no way to produce outputs with **multiple characters in a row**. Consider the alignment [h, h, e, l, l, l, o]. Collapsing repeats will produce "helo" instead of "hello".

Blank Token

ϵ

To get around these problems, CTC introduces a new token to the set of allowed outputs.

This new token is sometimes called the **blank token**. We'll refer to it here as ϵ .

The ϵ token doesn't correspond to anything and is simply removed from the output.

The alignments allowed by CTC are the **same length as the input**.

We allow **any alignment which maps to Y** after merging repeats and removing ϵ tokens



First, merge repeat characters.



Then, remove any ϵ tokens.



The remaining characters are the output.

h e | | o

If Y has two of the same character in a row, then a valid alignment **must have an ϵ** between them.

With this rule in place, we can differentiate between alignments which collapse to "hello" and those which collapse to "heloo".

Let's go back to the output $[c, a, t]$ with an input of length six. Here are a few more examples of valid and invalid alignments.

Valid Alignments



Invalid Alignments



corresponds to
 $Y = [c, c, a, t]$

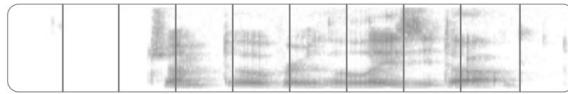
has length 5

missing the 'a'

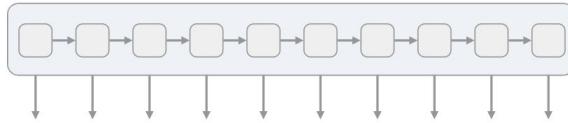
First, the allowed alignments between X and Y are monotonic. If we advance to the next input, we can keep the corresponding output the same or advance to the next one.

A second property is that the alignment of X to Y is many-to-one. One or more input elements can align to a single output element but not vice-versa.

This implies a third property: the length of Y cannot be greater than the length of X .



We start with an input sequence,
like a spectrogram of audio.



The input is fed into an RNN,
for example.

h	h	h	h	h	h	h	h	h	h
e	e	e	e	e	e	e	e	e	e
l	l	l	l	l	l	l	l	l	l
o	o	o	o	o	o	o	o	o	o
€	€	€	€	€	€	€	€	€	€

The network gives $p_t(a | X)$,
a distribution over the outputs
 $\{h, e, l, o, \epsilon\}$ for each input step.

h	e	€			€			o	o
h	h	e			€	€		€	o
€	e	€			€	€		o	o

With the per time-step output
distribution, we compute the
probability of different sequences:

h	e			o
e			o	
h	e		o	

By marginalizing over alignments,
we get a distribution over outputs

$$p(Y \mid X) = \sum_{A \in \mathcal{A}_{X,Y}} \prod_{t=1}^T p_t(a_t \mid X)$$

The CTC conditional
probability

marginalizes over
the set of valid
alignments

computing the **probability** for a
single alignment step-by-step.

CTC Score

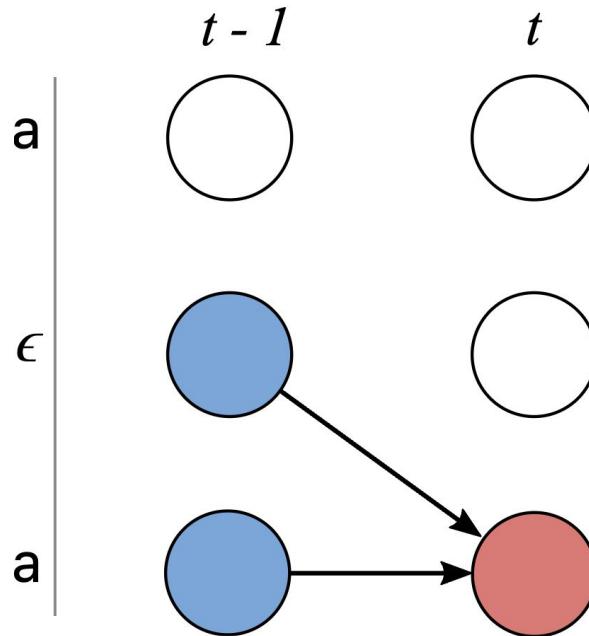
Since we can have an ϵ before or after any token in Y , it's easier to describe the algorithm using a sequence which includes them. We'll work with the sequence

$$Z = [\epsilon, y_1, \epsilon, y_2, \dots, \epsilon, y_U, \epsilon]$$

which is Y with an ϵ at the beginning, end, and between every character.

Let's let α be the score of the merged alignments at a given node. More precisely, $\alpha_{s,t}$ is the CTC score of the subsequence $Z_{1:s}$ after t input steps. As we'll see, we'll compute the final CTC score, $P(Y | X)$, from the α 's at the last time-step. As long as we know the values of α at the previous time-step, we can compute $\alpha_{s,t}$. There are two cases.

Case 1



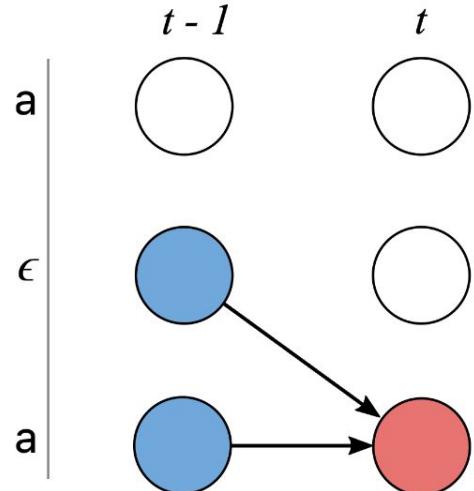
$$\alpha_{s,t} = (\alpha_{s-1,t-1} + \alpha_{s,t-1}) \cdot p_t(z_s \mid X)$$

The CTC probability of the two valid subsequences after $t - 1$ input steps.

The probability of the current character at input step t .

Case 1:

In this case, we can't jump over z_{s-1} , the previous token in Z . The first reason is that the previous token can be an element of Y , and we can't skip elements of Y . Since every element of Y in Z is followed by an ϵ , we can identify this when $z_s = \epsilon$. The second reason is that we must have an ϵ between repeat characters in Y . We can identify this when $z_s = z_{s-2}$.



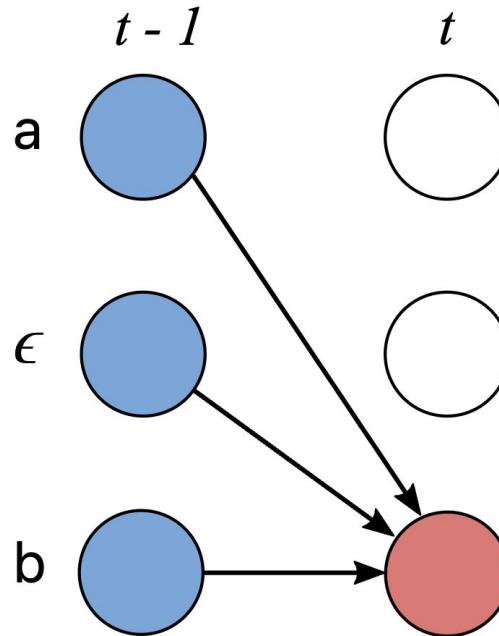
To ensure we don't skip z_{s-1} , we can either be there at the previous time-step or have already passed through at some earlier time-step. As a result there are two positions we can transition from.

$$\alpha_{s,t} = (\alpha_{s-1,t-1} + \alpha_{s,t-1}) \cdot p_t(z_s | X)$$

The CTC probability of the two valid subsequences after $t - 1$ input steps.

The probability of the current character at input step t .

Case 2



$$\alpha_{s,t} = (\alpha_{s-2,t-1} + \alpha_{s-1,t-1} + \alpha_{s,t-1}) \cdot$$

The CTC probability of the three valid subsequences after $t - 1$ input steps.

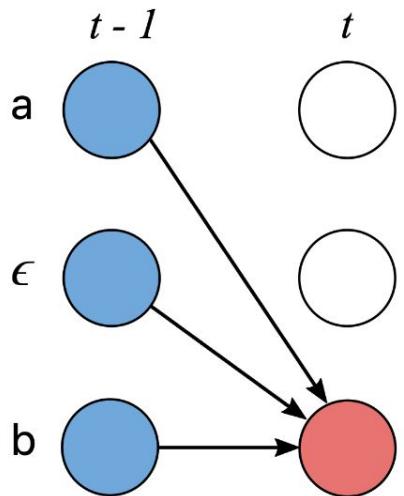
$$p_t(z_s \mid X)$$

The probability of the current character at input step t .

Case 2:

In the second case, we're allowed to skip the previous token in Z .

We have this case whenever z_{s-1} is an ϵ between unique characters. As a result there are three positions we could have come from at the previous step.

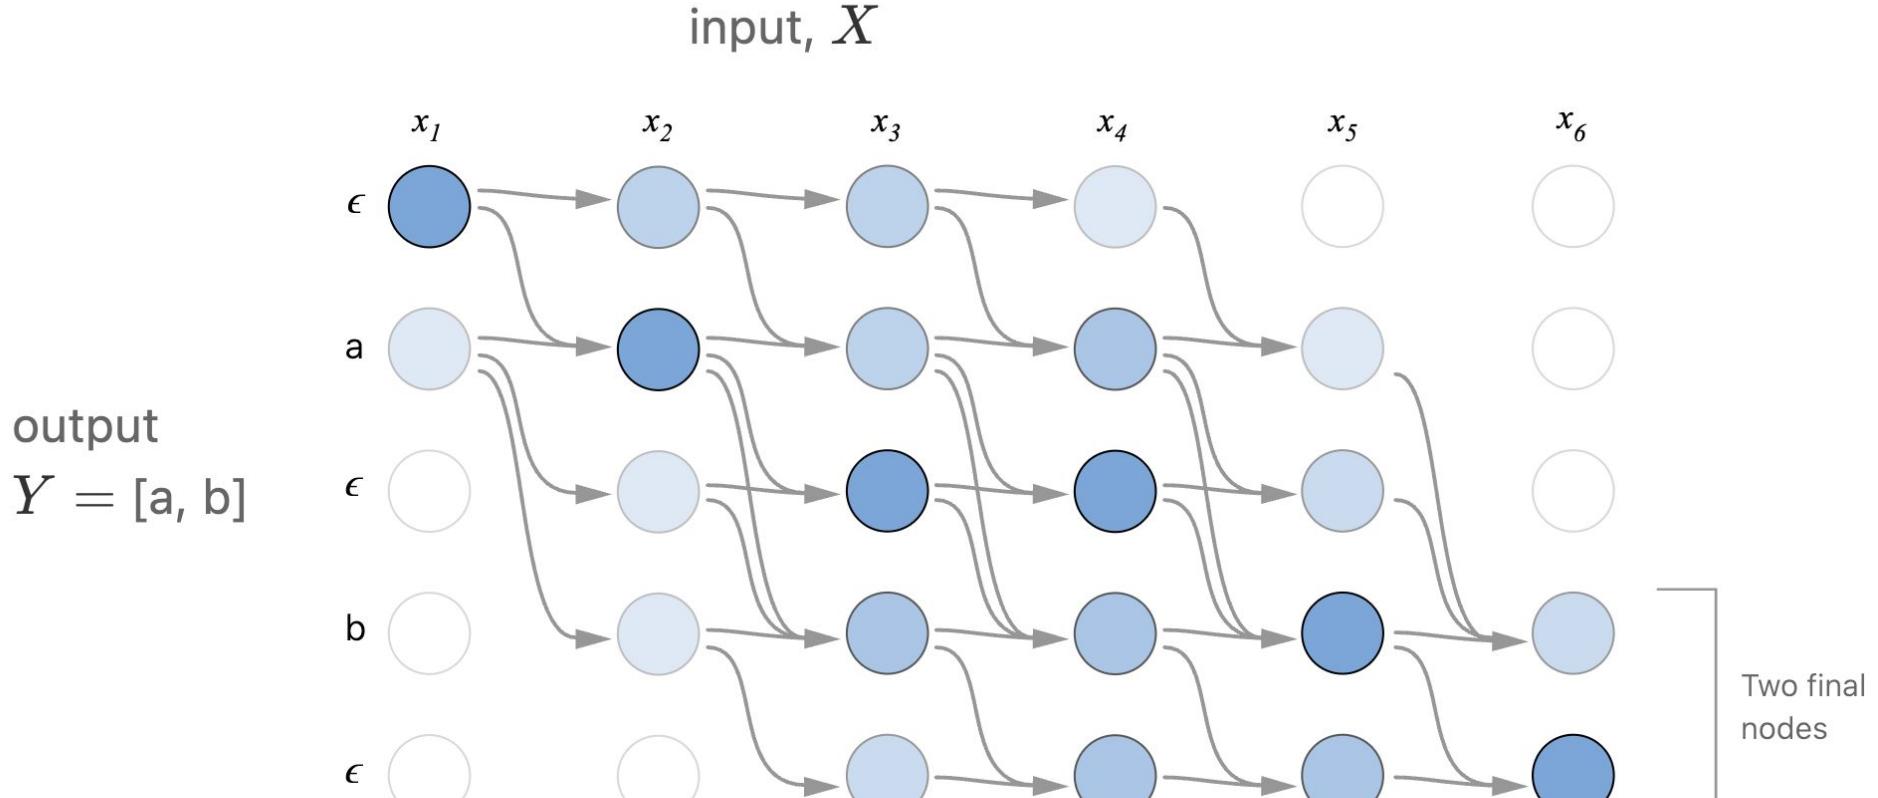


$$\alpha_{s,t} = (\alpha_{s-2,t-1} + \alpha_{s-1,t-1} + \alpha_{s,t-1}) \quad \cdot$$

The CTC probability of the three valid subsequences after $t - 1$ input steps.

$$p_t(z_s \mid X)$$

The probability of the current character at input step t .



Node (s, t) in the diagram represents $\alpha_{s,t}$ – the CTC score of the subsequence $Z_{1:s}$ after t input steps.

CTC Loss Function

For a training set \mathcal{D} , the model's parameters are tuned to minimize the negative log-likelihood

$$\sum_{(X,Y) \in \mathcal{D}} -\log p(Y | X)$$

instead of maximizing the likelihood directly.

Inference

Inference

After we've trained the model, we'd like to use it to find a likely output for a given input. More precisely, we need to solve:

$$Y^* = \operatorname{argmax}_Y p(Y | X)$$

One heuristic is to take the most likely output at each time-step. This gives us the alignment with the highest probability:

$$A^* = \operatorname{argmax}_A \prod_{t=1}^T p_t(a_t | X)$$

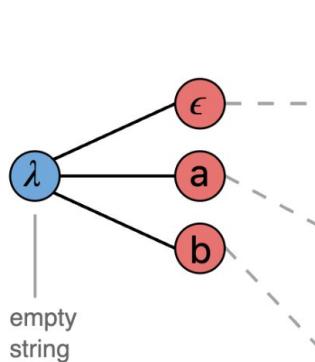
We can then collapse repeats and remove ϵ tokens to get Y .

For many applications this heuristic works well,
especially when most of the probability mass is allocated to a single alignment.
However, this approach can sometimes miss easy to find outputs with much higher probability.

A regular beam search computes a new set of hypotheses at each input step.

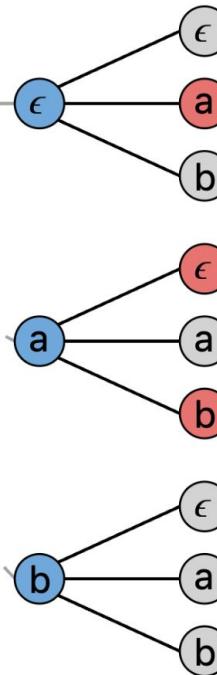
$T = 1$

current hypotheses proposed extensions



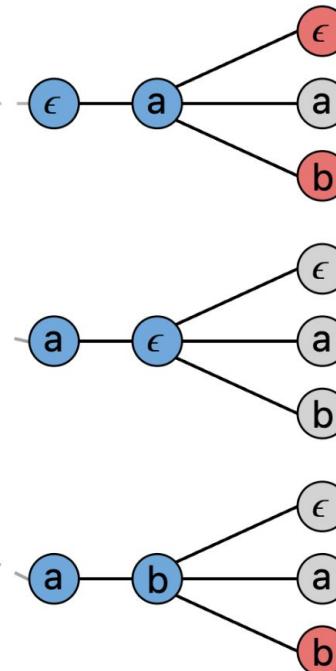
$T = 2$

current hypotheses proposed extensions



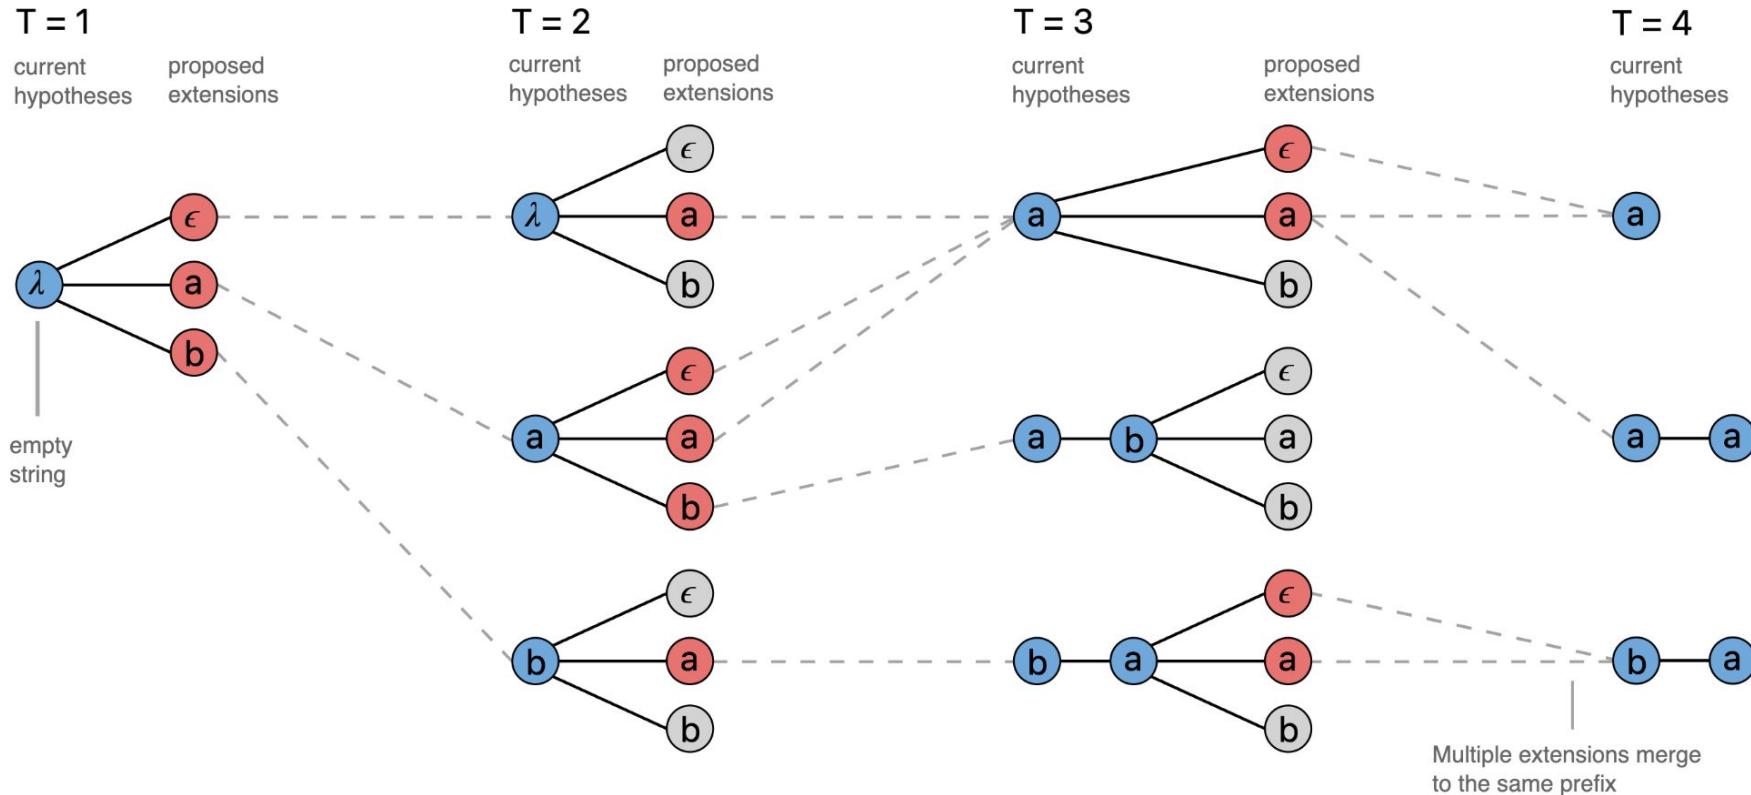
$T = 3$

current hypotheses proposed extensions



A standard beam search algorithm with an alphabet of $\{\epsilon, a, b\}$ and a beam size of three.

We can modify the vanilla beam search to handle multiple alignments mapping to the same output. In this case instead of keeping a list of alignments in the beam, we store the output prefixes **after collapsing repeats and removing ϵ characters**.



The CTC beam search algorithm with an output alphabet $\{\epsilon, a, b\}$ and a beam size of three.

We have to keep track of two probabilities for each prefix in the beam.

The probability of all alignments which end in ϵ and the probability of all alignments which don't end in ϵ .

